

ECE/CS 466 Project 1

Due: March 31, 2015

In this part of the project, you are asked to learn how to use the popular architectural simulator, SimpleScalar, to analyze the performance impact of varying architectural parameters.

First, you need to download “simplesim-3v0e.tgz” from www.simplescalar.com and install the simulator on your computer. If your computer is windows-based, you need to install the simulator either under “cygwin” or a Linux-like virtual machine.

When you install the codes, please use the PISA configuration (follow the installation instruction in file “README” and use “make config-pisa”). To verify that you have installed everything correctly, run the testing with command of “./sim-safe tests/bin.little/test-math”. You should see printed results as shown in the class.

Next, you will use SimpleScalar to run a SPEC2000 benchmark program, equake. This is a floating-point application for Seismic Wave Propagation Simulation. You can download the PISA binary code of equake (equake.ss) and the input file (equake.in) from the blackboard. It’s better to put the two files under the same directory of sim-outorder and sim-safe on your computer.

Now, you can begin the following experiments to see the performance impact of varying hardware parameters. Run the simulation and report the observed results (including the printed statistics in your report – only the related ones, not the full printout). Since running the program from beginning to end using the detailed simulator would take too long, you are asked to run only a small portion of the program. Please use the following option “-fastfwd 500000000” to fast-forward the first 500 million instructions and collect statistics on the next 300 million instructions (using the option “-max:inst 300000000”).

1. What is the performance of running the program, equake, under the default system setup (without changing any simulation parameters) using command: `./sim-outorder equake.ss < equake.in`?
2. How much is the performance loss if the processor uses in-order execution instead of the default out-of-order execution for running the program?
3. The above experiments only perform detailed simulation on 300 million instructions. Based on the simulator running time in Question 1, estimate how long it would take to run the program on a real machine with 3GHz processor and the same IPC value as in Question 1; and then estimate how long it would take to simulate the program’s execution in details from beginning to end using the default configuration. Note: Do not run the detailed simulation from beginning to end. It may take days to finish.

4. An advantage of using simulator is that it can help us to identify the performance bottleneck of a design. Start from the default configuration and double the pipeline width on fetch, decode, issue and commit stages at the same time (keeping other parts of the processor including FUs the same). What would be the performance improvement compared with the default configuration? Now since the pipeline width on each stage has been doubled, the default size of some hardware components such as FUs or load/store buffers might be too small for such configuration and limit the performance. Their sizes need to be increased too in order to have a balanced pipeline. Which hardware component needs to increase its size first? Use simulation results to support your claim.

Hint:

- You need to use “sim-outorder” in order to get any performance-related statistics such as the CPI or IPC value.
- There is no program called sim-inorder for in-order execution. But sim-outorder has an option to perform in-order execution.
- The reported “sim_elapsed_time” is how fast the simulator runs on your machine instead of how fast the program runs on the simulated processor.
- For question 3, you can use “sim-safe” to run the program till the end to get the total number of instructions of the program (it may take a few hours depending on the speed of your machine) and then do the estimation.